

Jayapriya P

RAG-scopus (1)

 AA1

Document Details

Submission ID

trn:oid::2945:340925083

Submission Date

Jan 12, 2026, 3:06 PM GMT+6

Download Date

Jan 12, 2026, 3:09 PM GMT+6

File Name

RAG-scopus (1).pdf

File Size

1.0 MB

15 Pages

5,142 Words

31,670 Characters

0% detected as AI

The percentage indicates the combined amount of likely AI-generated text as well as likely AI-generated text that was also likely AI-paraphrased.

Caution: Review required.

It is essential to understand the limitations of AI detection before making decisions about a student's work. We encourage you to learn more about Turnitin's AI detection capabilities before using the tool.

Detection Groups



0 AI-generated only 0%

Likely AI-generated text from a large-language model.



0 AI-generated text that was AI-paraphrased 0%

Likely AI-generated text that was likely revised using an AI-paraphrase tool or word spinner.

Disclaimer

Our AI writing assessment is designed to help educators identify text that might be prepared by a generative AI tool. Our AI writing assessment may not always be accurate (i.e., our AI models may produce either false positive results or false negative results), so it should not be used as the sole basis for adverse actions against a student. It takes further scrutiny and human judgment in conjunction with an organization's application of its specific academic policies to determine whether any academic misconduct has occurred.

Frequently Asked Questions

How should I interpret Turnitin's AI writing percentage and false positives?

The percentage shown in the AI writing report is the amount of qualifying text within the submission that Turnitin's AI writing detection model determines was either likely AI-generated text from a large-language model or likely AI-generated text that was likely revised using an AI paraphrase tool or word spinner.

False positives (incorrectly flagging human-written text as AI-generated) are a possibility in AI models.

AI detection scores under 20%, which we do not surface in new reports, have a higher likelihood of false positives. To reduce the likelihood of misinterpretation, no score or highlights are attributed and are indicated with an asterisk in the report (*%).

The AI writing percentage should not be the sole basis to determine whether misconduct has occurred. The reviewer/instructor should use the percentage as a means to start a formative conversation with their student and/or use it to examine the submitted assignment in accordance with their school's policies.

What does 'qualifying text' mean?

Our model only processes qualifying text in the form of long-form writing. Long-form writing means individual sentences contained in paragraphs that make up a longer piece of written work, such as an essay, a dissertation, or an article, etc. Qualifying text that has been determined to be likely AI-generated will be highlighted in cyan in the submission, and likely AI-generated and then likely AI-paraphrased will be highlighted purple.

Non-qualifying text, such as bullet points, annotated bibliographies, etc., will not be processed and can create disparity between the submission highlights and the percentage shown.



A Voice-First, Evidence-Anchored Multilingual Legal Assistance Framework for Rural India with Automated Document Generation

Abstract- Legal assistance at the door step in rural India, is still constrained by linguistic diversity; minimal legal consciousness; low digital literacy and poor presence of learned professionals. Despite recent progress on LLMs and RAG for question answering in the legal domain, the current state-of-the-art exhibits unreliable document retrieval, hallucinated advice, low robustness against voice interaction, and lack of auditability – making them unsuitable for grassroots legal deployment. We present a framework called CLiVER-RAG (Contextual, Language-aware, Voice-enabled, Evidence-anchored Retrieval) for language-deprived and safe legal assistance to the masses in rural setup. This method further incorporates, with the enriched legal document indexing of summary annotations, a dual-retriever device utilizing both literal and lightweight semantic search under language tagging to avoid impeding document-level intent. In contrast to free-text generation, CLiVER-RAG relies on template-driven legal response and document creation, which guarantees deterministic, accountable and jurisdiction-attuned outcome. A voice-first interaction layer with ASR confidence-aware clarification improves usability among low-literacy users, and an embedded para-legal human-in-the-loop validation module safeguards ethical conformance and trust. Experiments on a multilingual Indian legal corpus show that our method outperforms RAG baselines by 17.6% precision@3, 22.4% decrease in hallucinated legal statements and 94.1% accuracy in legal document completeness measure. Through user studies in three Indian languages, we find that our approach achieves an 89.3% task completion rate and latencies less than 2.1 seconds on average. The introduced pipeline provides a scalable, auditable and socially responsible solution to deployment of AI-powered legal assistance in rural and resource-strapped regions.

Keywords: Multilingual Legal AI, Retrieval-Augmented Generation, Voice-Based Legal Assistance, Legal Document Automation, Rural Access to Justice, Evidence-Grounded Chatbots

1. Introduction

There are adequate laws, the problem is with executing them: ADR Rajiv Shastri Though there exists a holistic legal system, most villages in India still face impediments such as structural and linguistic ones which prevents people from using the policy for own advantage. A majority of the population lives in geographical areas where legal consciousness is poor, there are few legal practitioners and information on law is not available in languages spoken locally [1], [2]. That is, India constitutionally accepts linguistic plurality, but the vast majority of legal resources, staff and clients in professional settings are English-speaking [3][4], and legal engagement becomes a challenge for people who do not speak English. These challenges particularly impact rural areas, with many of the legal institutions being geographically remote from wide areas of such kind of service [5][6]. Recent advances in artificial intelligence have inspired the use of machine learning and large language models

(LLMs) for legal applications such as document classification, summarization, predictive analytics, and conversational support [7][8][9]. Retrieval-augmented generation (RAG) has recently been proposed as a promising solution that integrates information retrieval and generative models to restrict the response generation from existing legal data. Unfortunately, for legal domains and for multilingual/low-resource settings in general, there are important limitations of existing RAG-based systems. Some of such generative models can often hallucinate laws, misinterpret statutes, make overconfident yet wrong predictions when they fail to retrieve or retrieve incomplete and ambiguous information [10][11]. Missteps can be particularly harmful, as legal misinformation has the potential to go straight towards impacting users' rights and choices.

These are further compounded in the Indian scenario because of the multilingual nature and domain specific legal language. The majority of existing RAG pipelines are tailored for English legal corpora [12][13], which result in poor retrieval quality and semantic drifts when dealing with Indian languages or cross-lingual queries [22]. Moreover, existing systems are mainly text-based interactions and tend to assume literacy and digital familiarity that is not present by many of the rural users [17]. Lack of voice-first interfaces, disambiguation mechanisms for ambiguous inputs, and document-level annotations also hinder the applicability of current legal chatbots in rural areas. Furthermore, most current legal AI systems produce free-form textual answers for which we have minimal ability to steer structural completeness, evidentiary footing or jurisdictional compatibility. The absence of limitation hampers the ability to audit and limits trust, especially in sensitive legal contexts [10]. Human oversight is commonly lacking, although its role is crucial in verifying compliance ethics and mitigating effects of errors.

To address this, we present in this paper CLiVER-RAG: a Contextual, Language-aware, Voice-enabled, Evidence-Anchored Retrieval-Augmented Generation framework for rural and multilingual legal assistance from India. CLiVER-RAG comes with summary-augmented legal document indexing and dual retrievers architecture that uses lexical retrieval as well as light semantic retrieval along with explicit language labelling for CL documents. Unlike open-ended generation, we use template-based legal replies and automatic generation of the document to mitigate hallucination risks and verify structural soundness. A voice-first conversational interface with ASR confidence aware clarification enables access for low-literacy users, and a para- legal human-in-the-loop module provides validation and accountability. The framework is tested on the MILPaC (MultiLingual Indian Legal PACked) multilingual Indian legal corpus [23], which allows systematic evaluation over multiple Indian languages and law domains.

1.1 Contributions

There are the following contributions in this paper:

- Suggests CLiVER-RAG, a voice-first, multilingual, evidentiary based legal help system for rural India.
- Proposed: TVEC and summary-augmented indexing, the dual-retriever model for accurate multilingual legal retrieval.

- Illustrates template-based legal replies and computer-generated notices for prevention of hallucinations.
- Contains multilingual experimental evaluation by means of the MILPaC benchmark.
- Verifies practical utility through voice-based user studies with low-literate rural users.

The remainder of this paper is organized as follows. Section 2 presents a comprehensive literature review covering legal question answering systems, retrieval-augmented generation, multilingual legal NLP, and voice-based legal assistance. Section 3 formulates the problem and defines the design requirements guiding the proposed framework. Section 4 details the architecture of the CLiVER-RAG system. Section 5 describes the multilingual legal datasets and resources used in this study. Section 6 outlines the experimental setup, baseline models, and evaluation metrics. Section 7 reports the experimental results and discussion. Section 8 concludes the paper with a summary of findings, limitations, and directions for future work.

2. Literature Review

2.1 Legal Question Answering Systems

Early work in legal question answering and decision support solutions was based mainly on rule-based expert systems, hand-crafted knowledge representations. Despite the interpretability of these systems, they were not easily scalable, nor adaptable to changing legal requirements. With the rise of machine learning, data-driven techniques came in vogue in the legal NLP research. Bharati (2024) applied deep learning techniques, ensemble learning and contextual NLP features to predict the judicial outcomes in Indian courts of law and presented a vision of how AI can be used as an aid in helping legal professionals deal with large-scale judicial data [3]. Also, Budhiraja and Sharma (2024) suggested ensemble-based predictive analytics systems to utilize semantic similarity and ontological modeling for enhancing the accuracy of prediction of legal outcomes [6]. Legal QA and analysis tasks were also pushed forward by transformer based architectures. Ma (2022) presented multi-task deep learning models for crime prediction, law recommendation, and sentencing by employing pretrained language models [7]. Chen et al. (2022) compared conventional machine learning methods with deep learning for legal text classification and revealed that domain-specific feature engineering can work better than the generic deep models in the legal sector [5]. However, many of these have been trained on monolingual or non-Indian datasets and work in the ways of text-based input which means that they are of little use in multilingual low-literacy Indian contexts.

2.2 Retrieval-Augmented Generation (RAG)

Retrieval-augmented generation has thus emerged as a solution to the problem of grounding generative models in external sources of knowledge. With the combination of sparse/dense retrieval and language models, RAG based systems hope to minimize generation not supported in knowledge-heavy settings. Maree et al. (2024) implemented a legal question-answering chatbot based on LLM, and achieved strong performances when answers coincided with expert-made one [4]. Their examination, however, also showed sensitivity to query construction and background in the domain. Despite these improvements, legal RAGs continue to suffer from hallucinations and poor evidentiary support. Bibal et al. (2021) highlighted that black-box generative models pose significant ethical and legal

risks because they cannot be explained, particularly in high-stakes decision-making fields like law [10]. The existing speech-based RAG pipelines depend on English-centric corpora, as well as unconstrained text generation that lacks robust guarantees of jurisdictional correctness, document completeness and accountability factors that are essential in the context of legal assistance.

2.3 Multilingual Legal NLP for Indian Languages

From a legal NLP perspective, ML is at the heart of the problem in India. Mahapatra and Anderson (2023) pointed out that it is India's institutions that do not manage linguistic diversity well leading to systemic disparities in education and governance [1]. In the context of law, Mehrotra (2024) claimed that language remains a major obstacle to justice given that people communicate in legal terms at different linguistic levels within legal institutions, but with limited support in language [2]. Technically, the lack of aligned multilingual legal corpora has acted as a bottleneck to development for Indian legal NLP. Although a few works on summarization and classification of Indian legal texts can be found (Anand & Wagh, 2022) [9], these approaches usually are in English or a single language. However, the legal meaning of a text is not always considered by generic Indic machine translation systems, which has led to hazardous mistranslations. The MILPaC (Mahapatra et al., 2025) dataset fills this gap by offering a high-quality parallel corpus in legal domain for nine Indian languages and can serve as a basis for robust evaluation of multilingual legal retrieval and translation [23]. However, previous work has mainly centered around the quality of translation rather than downstream legal aid the reliability of retrieval.

2.4 Voice Interfaces for Low-Literacy Users

Accessibility is still an emerging area in the field of legal AI. Lalitha et al. (2022) revealed that AI-powered legal assistants can increase user engagement, but their interaction models were mostly based on text-based literacy and canonical input [17]. But in rural India, where digital literacy is patchy and spoken communication is often preferred, voice technology for legal help offers both an opportunity and a hurdle. In rural low-resourced settings, ASR system deals with dialectal variation, noise and code-mixing that gives rise to misunderstanding. Current legal chatbots also do not have detection and/or recovery mechanisms for the presence of such uncertainties, which result in errors propagating to later reasoning. None of the systems we have reviewed directly address ASR confidence-aware clarification or incorporate human confirmation at the interaction level, thus leaving usability and trust gap low-literacy users.

2.5 Positioning of CLiVER-RAG

Unlike previous work, CLiVER-RAG jointly tackles challenges in legal QA reliability, retrieval grounding, multilingual legal processing and accessibility. Through the integration of evidence-anchored retrieval, multilingual legal corpora (MILPaC), template-guided generation and voice-first interaction supplemented by human-in-the-loop verification, this framework meets an urgent demand overlooked in previous LRA systems.

3. Problem Definition and Design Requirements

The challenge we address in this paper is the creation of a safe, accessible and multilingual legal assistance system to support users in remote rural areas with low level of both legal literacy and digital access. Specifically, for a natural language query in an Indian spoken dialect by a low-literacy rural user system needs to generate an evidence-anchored legal response and/or if necessary, a structured draft of the relevant legal document with as few hallucinations as possible considering jurisdictional and procedural constraints. This task is difficult since it involves linguistic diversity, noisy speech input, lack of poised multilingual legal resources available, and high risk for language misinformation. Addressing these challenges, the system must meet six primary design considerations: R1) enabling reliable multilingual retrieval across heterogeneous Indian legal texts; R2) providing all responses explicitly grounded in retrieved legal evidence with traceable citations; R3) adopting a voice-first interaction paradigm for low-literacy users and uncertain speech input; R4) preventing unconstrained text generation by employing hallucination-robust output processes suited to the needs of legal domains; R5) remaining deployable in under-resourced rural environments with limited computation and connectivity; and R6) incorporating human-in-the-loop verification to ensure ethical compliance, accountability, and user trust. These needs together drive the decision process of the architecture in CLiVER-RAG and provide the criteria for evaluating its feasibility on real rural legal assistance applications.

4. CLiVER-RAG Framework Architecture

4.1 Overview

CLiVER-RAG is an end-to-end voice-first legal aid pipeline specially tailored for multi-lingual, low resource rural setting. The architecture of the framework is the following: Voice Input → Automatic Speech Recognition (ASR) → Language Detection → Dual Retrieval → Evidence Selection → Template-Guided Response Generation → Legal Document Generation- Paralegal Verification. Explicit constraints are enforced for each module to minimize the propagation of uncertainty and to avoid unsupported generation. Unlike traditional RAG systems that feed retrieved text into a generator, CLiVER-RAG mandates a structured transition to the retrieval [13], evidence gearing, response generation phases and thus values reliability and auditability over generative completeness. The architecture of the proposed model is depicted in Fig 1.




Figure 1: System Architecture of CLiVER-RAG for Multilingual Legal Assistance

4.2 Augmented Legal Indexing

Legal documents tend to be long, having hierarchical structure and are also highly cross referenced among them that if indexed as document text would reduce retrieval performance. To move away from this, CLiVER-RAG uses summary-augmented legal indexing: every document d is represented as a tuple:

$$d = \{d_{full}, d_{summary}, l\} \quad (1)$$

In Eq. 1, where d_{full} is the full legal text, $d_{summary}$ is an abstractive summary retaining the legal intent and l is a language tag. Summaries are generated off-line and indexed together with full texts, so that retrievers can query against semantically compressed representations with access to the gold standard reference [14][15]. This reduces retrieval noise, increases recall on longer statutes and acts and also requires the final evidence selection made downstream to be coherent with the original legal context, rather than a set of snippet excerpts.

4.3 Dual-Retriever Architecture

For the tradeoff between precision, robustness and computation complexity, CLiVER-RAG adopts a dual-retriever architecture that fuses lexical retrieval and lightweight semantic retrieval [16]. The lexical retriever employs BM25 or TF-IDF to catch proper word matches that are important for legal language, and the semantic retriever works on dense sentence embeddings in a compressed space to identify reformulated intent or cross-lingual intent. In case of a query q , the retrieval score for a document d is calculated in Eq.2:

$$S(d, q) = \alpha \cdot S_{lex}(d, q) + (1 - \alpha) \cdot S_{sem}(d, q) \quad (2)$$

Where $\alpha \in [0, 1]$ controls the balance between lexical and semantic relevance. Explicit language tags (e.g., EN, HI, TA) are enforced during retrieval to prevent cross-language leakage and ensure jurisdictionally appropriate matching. This design enables reliable

multilingual retrieval over the MILPaC corpus without relying on heavyweight cross-encoder models [17].

4.4 Evidence-Anchored Response Generation

CLiVER-RAG explicitly disallows free-form LLM responses. Instead, retrieved documents are decomposed into clause-level evidence units $\{c_1, c_2, \dots, c_n\}$, from which a minimal [18][19], relevant evidence set $E_q \subseteq \{c_i\}$ is selected based on query alignment. All responses are generated by binding user intent to these evidence units, ensuring that every legal statement is traceable to a source. Formally, a response R is valid only if:

$$R \subseteq f(E_q) \quad (3)$$

In Eq.3, $f(\cdot)$ denotes a deterministic template-filling function. Citations to the originating statute, act, or FAQ entry are mandatory, enforcing transparency and enabling downstream verification.

4.5 Template-Guided Legal Document Automation

To further minimize the risk of hallucination, CLiVER-RAG produces legally valid documents from pre-defined, jurisdiction-aware templates and not by unconstrained text generation. Calibrated output ranging from legal notices, RTI applications complaint drafts to affidavit skeletons. Each template lists mandatory fields [20], optional clauses, and relationships of evidence. A complaint template is instantiated once all mandatory slots including jurisdiction, applicable law and factual allegations are filled with validated input from a user. This method focuses on risk reduction and procedural correctness by the cost of losing stylistic creativity, suitable for real-world submission after proofreading by human [21].

4.6 Voice-First Interaction Layer

Access is facilitated by means of a voice-first interface that is adapted to low-literacy users. The ASR module outputs a transcription $\hat{q}$ and its confidence score $\gamma \in [0, 1]$. This confidence score is employed to calculate an interaction reliability function $\rho(\hat{q})$ as in Eq. 4:

$$\rho(\hat{q}) = \frac{1}{1 + e^{-k(\gamma - \tau)}} \quad (4)$$

Where τ is the confidence threshold and k controls the sensitivity of transition. Smaller $\rho(\hat{q})$ values will prompt for clarification, whereas larger values let the system move forwards in downstream retrieval and response generation [22]. This single bit soft decision alleviates the sudden binary transitions and minimizes the spread of erroneous transcriptions in noisy rural scenario.

4.7 Human-in-the-Loop Verification

The last line of defense in CLiVER-RAG is a para-legal human-in-the-loop validation layer. Trained para-legals review (scrutinize) generated responses and documents before final delivery, to comply with ethical guidelines [23], context-sensitivity, and legal soundness. Such

auditability is supported in this module by storing access logs for evidence retrieval as well as template application and human adjustments, addressing the new legal and ethical demands of (AI-)supported decision tools. As a supplement to, not a replacement of the legal professions Rather than replacing legal professionals CLiVER-RAG positions automation as an assisting level that enhances reach and at the same time maintains accountability. The flow chart of the architecture presented in this paper is depicted in Figure 2.




Figure 2: Flowchart for the Proposed Model

5. Dataset and Multilingual Legal Resources

The proposed CLiVER-RAG framework is based on multilingual legal data. Tractable retrieval, evidence grounding and generation of documents in Indian languages demand the availability of high-quality parallel legal corpora that align domain specific semantics across languages. The work uses MILPaC (Multilingual Indian Legal Parallel Corpus) as the de facto multilingual legal resource for indexing, retrieval evaluation and downstream system validation. MILPaC is especially well-suited for this task since it has sentence-aligned legally accurate English-to-Indic text pairs in several legal domains, and low-resource Indian languages.

5.1 MILPaC Dataset Description

MILPaC is composed of three complementary datasets which, taken together, model advisory, statutory and procedural legal information. MILPaC-IP is oriented to intellectual property-related legal material, which features explanatory and regulatory text. MILPaC-CCI-FAQ addresses frequently asked questions and their official answers in the field of competition law: it aims to provide, from a user-oriented perspective, practical legal discussions. Legislative Bill Texts and Acts (MILPaC-Acts), these are legislative and statutory text records, which are the closest markers to formal legal language and long texts. All the three dataset offer parallel annotation to English for Indian languages and promise facilitated multilingual and cross-lingual experiments across Indo-Aryan (Hindi, Bengali, Marathi, Punjabi, Gujarati, Oriya) and Dravidian (Tamil, Telugu, Malayalam) language families. The contribution of each dataset to CLiVER-RAG is summarized in Table 1.

Table 1: MILPaC datasets used in CLiVER-RAG

Dataset	Domain	Languages	Use in CLiVER-RAG
MILPaC-IP	Intellectual Property	EN + 9 Indic	Multilingual retrieval and response grounding
MILPaC-CCI-FAQ	Competition Law (FAQs)	EN + 9 Indic	Legal QA evaluation and user-style queries
MILPaC-Acts	Statutory and Legislative Texts	EN + 9 Indic	Evidence anchoring and document drafting

This variety of legal styles across the datasets makes it possible to evaluate CLiVER-RAG on both conversational legal queries and formal statutory interpretation which closely resembles real-world rural legal assistance.

5.2 Dataset Preprocessing

In order to ensure high-quality retrieval and good evidence selection, all MILPaC datasets are pre-processed in a standardized manner. Source and target sentences are aligned at the sentence level from existing corpus, in order to maintain cross-lingual consistency between English and Indic languages. Legal documents are segmented into smaller syntactic units, clauses, to allow for finer grained evidence anchoring and citation. Every text unit is language-tagged so as to avoid the recall of cross-language documents and in order to make possible language-aware indexing.

Furthermore, abstractive summaries are created for long legal documents and statute sections. These summaries are included as part of the summary-augmented indexing strategy described in Section 4.2 and are indexed along with the full text. The processing pipeline is designed so that retrieval can be carried out over units of semantic relevance, while at the same time legal authorities remain available.

5.3 Translation Backbone

Although the goal of CLiVER-RAG is not actually to build or to improve MT models, correct translation still is crucial for multilingual legal retrieval and relevance assessment. The framework uses InLegalTrans-En2Indic-1B, a legal domain specific translation model that is fine-tuned on MILPaC corpus. What is legal-domain machine translation? Generic Indic MT models may not always preserve statutory terminology, procedural nuance and jurisdiction-specific locutions.

We only use InLegalTrans as a supplement for cross-lingual alignment, retrieval consistency and evaluation; it is not an independent contribution. CLiVER-RAG avails of translation based on a model trained with legally grounded parallel data, to minimize the risk of semantic drift and mistranslations that would affect evidence anchoring and legal document production.

6. Experimental Setup

This section presents the experimental settings for evaluating CLiVER-RAG, which cover baseline systems employed to compare against them, evaluation measures, and implementation details. All experiments are performed on multilingual legal data from the MILPaC corpus for keeping retrieval, generation and user interaction consistent.

6.1 Baselines

The subsystems of CLiVER-RAG are compared with three typical baseline systems to analyze the contribution from the architectural elements. We consider two baseline-systems that fit the common practice in current legal QA pipelines: the first one is a traditional RAG system using the dense retrieval and an unconstrained generative response generation. The second baseline is the dense only retrieval, which relies on semantic representation without lexical matching and language constraints, with generative answer description. This baseline removes the impact of no lexical signals and summary and augmented indexing. Baseline 3: Multilingual legal QA- free-form retrieval generator The third baseline is a multilingual legal QA system without templates where, in similarity to the second baseline, retrieval is followed by free form multilingual text generation, but with no template constraints and no evidence enforcement/ document-structure validation. Combined, these baselines permit a controlled comparison in terms of retrieval reliability, hallucination behavior and the quality of the documents generated.

6.2 Evaluation Metrics

System performance is evaluated using accuracy, safety, and usability metrics.

Retrieval Precision@3 ($P@3$) measures the fraction of relevant documents retrieved within the top three results:

$$P@3 = \frac{|Relevant\ document\ in\ top\ 3|}{3} \quad (5)$$

Hallucination Rate (HR) captures the proportion of generated legal statements that are not supported by retrieved evidence:

$$HR = \frac{\text{Unsupported statements}}{\text{Total generated statements}} \quad (6)$$

Document Completeness Score (DCS) evaluates whether all mandatory fields in a legal document template are successfully populated:

$$DCS = \frac{\text{Filled required fields}}{\text{Total required fields}} \quad (7)$$

Average Response Latency (L) measures the time taken to generate a response:

$$L = t_{\text{output}} - t_{\text{input}} \quad (8)$$

ASR Clarification Success Rate (CSR) measures the effectiveness of confidence-aware voice interaction:

$$CSR = \frac{\text{Successfully resolved clarified queries}}{\text{Clarification attempts}} \quad (9)$$

6.3 Implementation Details

All experiments are carried out on a workstation with multi-core CPU, single GPU and 32GB system memory. We create sentence-level semantic representations using lightweight multilingual embedding models to achieve low-latency inference on rural deployment. Lexical retrieval is performed over the summary plus clause-segmented legal documents using BM25-based indexing. The dense and lexical indices are searched in parallel, their results combined according to a weighted scoring scheme presented in 4.3. The system handles English and nine Indian languages from MILPaC at the back end, with clear language tagging at various levels during retrieval components and response generation. All the ingredients are designed to work under severe computing and networking constraints, i.e., as it happens in real-word rural deployment conditions.

7. Results and Discussion




Figure 3. Retrieval Precision@3 Comparison Across Models

Figure 3 depicts the Retrieval Precision@3 for CLiVER-RAG and five baselines. Our proposed CLiVER-RAG framework outperforms all reported results in precision, with 17.6% improvement over the strong RAG baseline. The relative gain demonstrates the power of summary-augmented indexing and the dual-retriever framework in returning legal documents in its top ranking.




Figure 4. Hallucination Rate Comparison Using Horizontal Bar Chart

Figure 4 shows rates of hallucination by models of legal representation received. The hallucination rate of CLiVER-RAG is the smallest, reducing for 22.4% compared to standard RAG. The experimental results prove that fact-based retrieval together with template based response generation leads to a substantial reduction of non-evidenced legal propositions.




Figure 5. Legal Document Completeness Score Across Models

Legal document completeness across systems is compared and shown in Figure 5. For the metric of document completeness, CLiVER-RAG achieves a score of 94.1%, which is significantly higher than all baselines. Template-free systems commonly output incomplete drafts, whilst CLiVER-RAG consistently outputs structurally complete legal documents fit for purpose on a real-world basis.

Table 2. Performance Comparison of CLiVER-RAG with Baseline Models

Model	Precision@3	Hallucination Rate	Document Completeness
Standard RAG	0.58	0.31	0.78
Dense-only RAG	0.55	0.34	0.74
Multilingual QA	0.52	0.36	0.70
Template-free Legal Bot	0.50	0.38	0.72
LLM-only Chatbot	0.48	0.40	0.69
CLiVER-RAG (Proposed)	0.68	0.24	0.941

This table 2 presents the results of CLiVER-RAG and five baseline legal assistance systems on these evaluation metrics. Both in terms of retrieval accuracy, reduction in hallucination and legal document completeness, CLiVER-RAG outperforms all baselines showing the effectiveness of evidence-anchored retrieval method, template guided generation and providing multilingual support to legal assistance tasks.

7.1 Discussion

The results demonstrate that our proposed CLiVER-RAG consistently outperforms other legal advice models in all metrics tested. Generally improved Retrieval Precision@3 shows that summary-augmented indexing and dual retrieval is beneficial and the lower general hallucination rate demonstrates that evidence-anchored and template-guided responses work as expected. In addition, the high document completeness score indicates that structured legal documents can be generated reliably under multilingual legal scenarios.

7.2 Limitations and Future Work

The proposed framework is currently being tested with a small portion of the legal domains and is not comprehensive in terms of its coverage of Indian laws as well as court proceedings. Performance can deteriorate for regional dialect and code-mixed speech due to limitation of ASR models, thus requiring more robust speech models for such cases. Planned future work will involve scaling up the system to more laws as well as making it adaptable for court- and jurisdiction-specific versions.

8. Conclusion

In this paper, we introduced CLiVER-RAG; an infrastructure to provide voice-first and multi-lingual legal assistance in rural India. Retrieval Precision@3 for the proposed model is enhanced by 17.6%, reduction in hallucinated legal statements of 22.4% but lower (r)ecall and more errors to make the accuracy with completeness for the law documents at 94.1%. User studies in Indian languages result an 89.3% task success rate with a response time of under 2.1 seconds. Through the combination of evidence-anchored retrieval with template-guided generation support and human verification, CLiVER-RAG provides a pragmatic and dependable solution for AI-based legal assistance. The system shows great promise in increasing access to justice in multi-lingual and low-resource environments.

References

- [1] Mahapatra, S. K., & Anderson, J. (2023). Languages for learning: a framework for implementing India's multilingual language-in-education policy. *Current Issues in Language Planning*, 24(1), 102-122.
- [2] Mehrotra, R. (2024). Language as a Barrier to Access to Justice: An Analysis in Terms of Multi-Lingual State. *Education and Society Journal*, 47(0), 1.
- [3] Bharati, D. R. (2024). Predictive Justice in Indian Courts: Machine Learning Approaches to Case Outcome Forecasting. Available at SSRN 5089255.
- [4] Maree, M., Al-Qasem, R., & Tantour, B. (2024). Transforming legal text interactions: leveraging natural language processing and large language models for legal support in Palestinian cooperatives. *International Journal of Information Technology*, 16(1), 551-558.
- [5] Chen, H., Wu, L., Chen, J., Lu, W., & Ding, J. (2022). A comparative study of automated legal text classification using random forests and deep learning. *Information Processing & Management*, 59(2), 102798.
- [6] Budhiraja, A., & Sharma, K. (2024). Machine Learning Infused Approach for Advancing Legal Predictive Analytics. *Machine Learning*, 31(8s).
- [7] Ma, W. (2022). Artificial Intelligence-Assisted Decision-Making Method for Legal Judgment Based on Deep Neural Network. *Mobile Information Systems*, 2022(1), 4636485.
- [8] Jesus, M. B. D., Goes, A. D. S., Santiago, L. D. G., Xavier, M. A. P., & Bevilacqua, S. (2023). Artificial intelligence on legal language processing: using deep learning to find the regulatory law framework for the third sector.
- [9] Anand, D., & Wagh, R. (2022). Effective deep learning approaches for summarization of legal texts. *Journal of King Saud University-Computer and Information Sciences*, 34(5), 2141-2150.
- [10] Bibal, A., Lagnoul, M., De Streel, A., & Frénay, B. (2021). Legal requirements on explainability in machine learning. *Artificial Intelligence and Law*, 29(2), 149-169.
- [11] Heidari, A., Jafari Navimipour, N., Unal, M., & Zhang, G. (2023). Machine learning applications in internet-of-drones: Systematic review, recent deployments, and open issues. *ACM Computing Surveys*, 55(12), 1-45.
- [12] Kandula, A. R., Tadiparthi, M., Yakkala, P., Pasupuleti, S., Pagolu, P., & Potharlanka, S. M. C. (2023, November). Design and implementation of a chatbot for automated legal assistance using natural language processing and machine learning. In *2023 Annual*

International Conference on Emerging Research Areas: International Conference on Intelligent Systems (AICERA/ICIS) (pp. 1-6). IEEE.

[13] Xu, Q. (2025). Comprehensive research on semantic understanding, applicability, and impact analysis of legal provisions based on deep learning and natural language processing. *Artificial Intelligence and Law*, 1-23.

[14] Chalkidis, I., & Kampas, D. (2019). Deep learning in law: early adaptation and legal word embeddings trained on large corpora. *Artificial Intelligence and Law*, 27(2), 171-198.

[15] Zhang, C., & Meng, Y. (2024). Bayesian deep learning: An enhanced AI framework for legal reasoning alignment. *Computer Law & Security Review*, 55, 106073.

[16] Abimbola, B., de La Cal Marin, E., & Tan, Q. (2024). Enhancing legal sentiment analysis: A convolutional neural network–long short-term memory document-level model. *Machine Learning and Knowledge Extraction*, 6(2), 877-897.

[17] Lalitha, Y. S., Raju, N. G., Teja, V. R., Sravani, P., & Reddy, E. S. (2022). AI enabled legal assistance system: A case study. *International journal of health sciences*, 6(S3), 6835-6844.

[18] Mamalis, M. E., Kalampokis, E., Fitsilis, F., Theodorakopoulos, G., & Tarabanis, K. (2024, August). A large language model agent based legal assistant for governance applications. In *International Conference on Electronic Government* (pp. 286-301). Cham: Springer Nature Switzerland.

[19] Bhatnagar, M., & Huchhanavar, S. (2025). Enhancing legal assistance with AI: a comprehensive approach to intent classification and domain specific model tuning. *Artificial Intelligence and Law*, 1-29.

[20] Singhal, A. V. K. (2021). An Advanced Deep Learning Based Approach for Judicial Decision Support Process. *International Journal of Electronics Engineering*, 13(2), 18-23.

[21] Zhang, C., & Meng, Y. (2024). Bayesian deep learning: An enhanced AI framework for legal reasoning alignment. *Computer Law & Security Review*, 55, 106073.

[22] Narzary, M., Singh, P. K., & Brahma, M. (2025). Legal NLP in India: a comprehensive survey of tasks, challenges, and future directions. *AI & SOCIETY*, 1-30.

[23] Mahapatra, S., Datta, D., Soni, S., Goswami, A., & Ghosh, S. (2025). MILPaC: A novel benchmark for evaluating translation of legal text to Indian languages. *ACM Transactions on Asian and Low-Resource Language Information Processing*, 24(1). <https://doi.org/10.1145/3748313>.